

SECTION 1

DEFINE THE PROBLEM

1. WHY

What business outcome are we trying to improve, and what does the problem cost today?

Time, money, errors, or churn. An estimate is fine. Is it more than the system costs to build, run, and be wrong?

What is the metric? Which candidate metric did we reject, and why?

Same goal, different metrics, different systems. The metric you pick is the system you get.

2. WHAT

What are we proposing to build? Is that a problem or a solution?

If you named a tool (chatbot, RAG, agent), write the problem underneath it instead.

What are we explicitly NOT building?

Three non-goals. An unstated non-goal is an invitation the system will accept.

3. WHO + ATOMIC UNIT

Who will use it, and who owns the decision it supports?

The user acts on the output. The owner's budget moves if it works. Name both. Then list who else cares; different seats often want different metrics.

What does the system decide about, one at a time?

The atomic unit is one _____

Examples: one customer-month, one support ticket, one passage, one task step.

4. HOW	How will it work? What simpler approach should we compare it against?
5. TRADE-OFFS	What are we optimizing for, and what are we willing to give up?
6. ASSUMPTIONS	What has to be true for this approach to work? <i>Include the signal: does the information this decision needs exist, and is it available at decision time?</i>

7. SIGNALS

What early evidence would tell us we are on track or off track?

Start from the top three ways this fails in production. Off-track signals come from failure modes.

On track

Off track

8. GOOD ENOUGH

What is good enough to test? Good enough to ship? What number makes us stop?

Pre-commit the stop with the owner; a number and a time horizon, agreed before building.

Test	Ship	Stop

Can you explain what this is for using no AI words at all?