

Signals Canvas

Fill this out with the person who owns the outcome, before you build. You're agreeing now on what working, failing, and an early warning look like, so the conversation later is *"the signal we agreed on has fired"* instead of *"I think we should stop."*

The goal

We're helping (<i>owner</i>)	
make this decision better	
using this approach	

Example: support agents, deciding what to tell a customer about a policy, using retrieval over the policy docs with a cited answer an agent reviews before sending.

Success: what working looks like

	Signal	Target	By when
Quantitative			
Qualitative			
Guardrail (<i>must not get worse</i>)			

Example: answer accuracy above 80% on the weekly production sample by day 60. Agents stop pinging the policy team with routine questions. Guardrail: escalations to the policy team fall rather than rise.

Stop: when we change course

	Signal	Threshold	By when
Quantitative			
Qualitative			
Fallback (<i>what we switch to</i>)			

Example: accuracy under 70% after two tuning cycles. Agents route around the bot and go back to asking the policy team. Fallback: keyword search plus an FAQ template. A stop signal without a named fallback is just a threat.

Leading: the early warning

	Signal	Threshold	Check on
Week one or two			
Week one or two			

Example: retrieval recall on a 50-question sample in week one. If the right policy isn't being found, accuracy can't reach 80% later.

Agreed by

Owner (name and role)	Date